

Zero Knowledge for Flock’s BLAKE3 Prover:

Construction, simulation, and concrete bounds

Alexander Lee

v7.1, July 2026 — implementation and security note

Abstract

Flock’s original BLAKE3 prover [7] was not zero knowledge: its transcript exposed deterministic functions of the witness. This paper gives a zero-knowledge variant for exactly two registered production profiles: a batch of 256 BLAKE3 instances and a batch of 256 fixed-digest statements

$$y_i = \text{BLAKE3}(x_i), \quad |x_i| = 64 \text{ bytes.}$$

The verifier is unchanged.

The proof has two algebraic steps. First, a field-valued mask makes the PIOP messages independent of the witness after conditioning on the two mask values that are later opened. For the production shape, failure occurs on the zero set of a nonzero polynomial of degree at most 720. Second, an affine translation of the PCS blinders preserves every opened value and every recursive codeword. These two translations reduce simulation to choosing the Fiat–Shamir challenges first, solving the dependent messages, and programming at most 18 fresh random-oracle points. The simulator receives only the public statement, and its output is accepted by the production verifier.

In the classical programmable random-oracle model, the explicit algebraic and programming terms sum to less than $2^{-118.502}$ at a hash-query budget $Q_H = 2^{64}$, before the ordinary 256-bit oracle-collision term. Thus the two registered profiles are computationally zero knowledge under the assumptions stated here. Knowledge security is separate: the available standalone Fiat–Shamir reduction over $\mathbb{F}_{2^{128}}$ gives 55.994 bits at the same query budget. The stronger 100-bit deployment target is conjectural, not a theorem of this paper.

1 Result and scope

Result in one sentence. For the two parameter sets in Table 1, a proof generated from a valid witness is computationally indistinguishable from a proof generated from the public statement alone, in the classical programmable-random-oracle model.

The fixed-digest relation is

$$\mathcal{R}_{\text{fd}} = \{((y_1, \dots, y_{256}), (x_1, \dots, x_{256})) : y_i = \text{BLAKE3}(x_i), |x_i| = 64 \text{ bytes}\}.$$

The batch relation uses the same BLAKE3 circuit but has no externally fixed digest list. In both cases the R1CS shape, packing convention, PCS parameters, and circuit digest are part of the registered statement family.

Throughout the paper, *computational zero knowledge* has its standard classical-ROM meaning: for every classical probabilistic polynomial-time distinguisher making at most Q_H oracle queries, the joint view of an honest proof and its oracle answers is computationally indistinguishable from

the view produced by a simulator that knows only the statement. The simulator may program points that the distinguisher has not already queried. This is the IOP and Fiat–Shamir setting of Ben-Sasson, Chiesa, and Spooner [1]; the algebraic masking step is in the spirit of zero-knowledge sumcheck [2].

Table 1: The complete production scope of the ZK claim.

Family	Public input	Circuit / PCS	Status
BLAKE3 batch	batch size 256	$m = 22$, rate 1/2, batch log 6	computational ZK
BLAKE3 preimage	256 digests	$m = 22$, rate 1/2, batch log 6	computational ZK

The scope is deliberately narrow. The preimage circuit hashes one message of exactly 64 bytes and pins the BLAKE3 IV, counter, block length, and `CHUNK_START`, `CHUNK_END`, and `ROOT` flags. Unknown circuit digests, batch sizes, and PCS parameters are rejected before proving. SHA-256 statements, Keccak statements, longer BLAKE3 messages, hash chains, Merkle-path statements, recursion, side channels, and quantum random-oracle queries are not covered.

The argument has a simple structure:

1. translate the PIOP masks while holding their revealed values fixed;
2. translate the PCS blinders while holding every opening fixed;
3. sample the remaining transcript from the public statement and program only fresh, domain-separated oracle points.

Sections 4–6 prove these steps in that order. Extraction is discussed separately because witness hiding and knowledge soundness are different properties.

Assumption 1 (Classical programmable random oracle). All injectively framed SHA-256 calls are modelled as one random function $\mathcal{H} : \{0, 1\}^* \rightarrow \{0, 1\}^{256}$. The adversary makes classical queries, and the simulator may program a point only before that point is queried. No QROM claim is made.

2 Relation and construction

Let z be the packed binary assignment for the block-diagonal BLAKE3 R1CS, embedded in $\mathbb{F}_{2^{128}}$. The prover establishes

$$(Az) \circ (Bz) = Cz$$

using Flock’s zerocheck and lincheck, and opens the committed vector with the ring-switched Ligerito PCS [3, 4]. For the fixed-digest family, the verifier derives an additional packed opening claim from the public digest list. Before the first challenge, the transcript binds that list, the padding rule, the witness layout, the circuit matrices, the parameter set, and a fresh 32-byte proof nonce.

The ZK prover adds four independent mask sources: the zerocheck mask P , the lincheck mask S , and the round-one masks S_c, S_h . The witness and mask commitments use independent hiding randomness, denoted μ and g . Every stream, including the proof nonce, is derived from a per-proof DRBG under a distinct label.

The main change is the zerocheck mask

$$\gamma P(X)Q^*(X), \quad Q^*(X) = q_0 + \sum_j q_j X_j.$$

Here P is uniform over $\mathbb{F}_{2^{128}}$ on a translated subcube of the top $d_{\log} = 12$ variables. The affine polynomial Q^* is public, protocol-fixed, and nonzero on the Boolean cube. The field-valued mask is essential: the withdrawn Boolean P, Q construction did not support the Schwartz–Zippel argument needed for production parameters. Protocol `flock-zk-fv-v3` therefore has no Q commitment, Q opening, final Q evaluation, or Q -randomness stream.

3 One random-oracle boundary

Fiat–Shamir, Merkle commitments, and proof of work use one oracle, but never the same input. Each input begins with a role byte. A Merkle input also binds the protocol version, proof nonce, commitment channel, recursion depth, node level, node index, payload length, and payload. A proof-of-work input is

$$\text{ROLE_POW} \parallel \text{state} \parallel \text{nonce}.$$

The proof nonce is absorbed immediately after the statement and before the first challenge.

This framing is part of the proof, not merely an encoding convention. It lets the extractor identify witness-leaf queries without relying on their order, prevents simulator programming from aliasing transcript or Merkle inputs, and makes proofs from different nonces or commitment channels use different oracle points. Native, recording, and programming backends implement the same frames. When no point is programmed they return identical bytes. Scalar and four-way SIMD implementations are differentially tested against the one-shot definition for every supported tail shape.

4 The PIOP translation

Fix the post-commitment challenge tuple c . Let D be the coefficient space of P . If $n = m - 6$ is the number of masked zerocheck rounds, define

$$R_c : D \longrightarrow \mathbb{F}_{2^{128}}^{2n}$$

to be the map from mask coefficients to the $2n$ round-message coordinates. Let

$$L_c : D \longrightarrow \mathbb{F}_{2^{128}}^2$$

map the same coefficients to the two mask values revealed later: the initial masked claim and $P(\rho)$. Once the leakage $L_c(P)$ is fixed, the remaining mask freedom is exactly $\ker L_c$, and the available change in the round messages is $R_c(\ker L_c)$.

Theorem 1 (Conditioned field-mask coverage). *For $m \in \{13, 15, 22\}$, the symbolic certificate establishes, outside the zero set of its selected minor,*

$$\text{rank} \begin{bmatrix} R_c \\ L_c \end{bmatrix} = 2n, \quad \text{rank } L_c = 2, \quad \dim R_c(\ker L_c) = 2n - 2.$$

For the production shape $m = 22$, the selected determinant is a nonzero polynomial of total degree at most 720 in the 32 challenge variables. Hence

$$\Pr_c[\text{coverage failure}] \leq \frac{720}{2^{128}} < 2^{-118.508}.$$

Proof sketch. Let $N = \dim D$. The stacked-rank statement gives

$$\dim(\ker R_c \cap \ker L_c) = N - 2n,$$

while $\text{rank } L_c = 2$ gives $\dim \ker L_c = N - 2$. Applying rank-nullity to R_c restricted to $\ker L_c$ yields

$$\dim R_c(\ker L_c) = (N - 2) - (N - 2n) = 2n - 2.$$

The implementation selects a $2n \times 2n$ minor and evaluates its determinant to a nonzero field element at one explicit challenge tuple. Its determinant is therefore not the zero polynomial. The generic symbolic kernels bound its total degree by 720, so Schwartz–Zippel gives the stated bad-set probability [5, 6]. $\square$

The concrete and symbolic matrices come from the same generic round kernels. The symbolic scalar has three interpretations: evaluation in $\mathbb{F}_{2^{128}}$, conservative multivariate degree bounds, and exact sparse polynomials for small instances. It has no inversion operation, so a challenge-dependent denominator cannot enter the calculation unnoticed. The selected rows, columns, nonzero evaluation, and degree vector are pinned in `s2_mask_coverage.json` and replayed by the certificate suite.

Theorem 1 proves the PIOP masking step. It says nothing by itself about privacy at the PCS opening boundary; that is the next step.

5 The PCS translation

Write the initial combined opening vector as

$$F = (\mu \parallel z) + c g,$$

where $c \in \mathbb{F}_{2^{128}}$ is sampled after commitment. Let d be a witness difference that preserves every public opening claim. The aim is to replace z by $z + d$ without changing anything the verifier opens.

Lemma 1 (Opening-preserving translation). *Suppose $c \neq 0$ and the initial-level query system has full row rank. There is an affine bijection on the uniform PCS masks such that replacing z by $z + d$ leaves the queried initial rows and the entire combined vector F unchanged. Consequently every recursive codeword, root, authentication path, sumcheck value, and out-of-domain sample is unchanged.*

Proof. Set

$$\delta g_{\text{top}} = c^{-1} d.$$

Choose $\delta \mu$ to solve the queried initial-level encoding equations, and set

$$\delta g_{\text{lo}} = c^{-1} \delta \mu.$$

Because the field has characteristic two, $d + c \delta g_{\text{top}} = 0$ and $\delta \mu + c \delta g_{\text{lo}} = 0$. Thus $\delta F = 0$ globally. The defining equations for $\delta \mu$ also hold the separately opened initial rows fixed. Translation by these deltas is a bijection on uniform masks, so it preserves their distribution. $\square$

The production opening queries 218 initial-level positions; each lane contains 512 independent mask symbols. After conditioning on all opened positions, each additional fresh wide leaf retains at least 16,384 bits of conditional entropy. No recursive entropy charge is needed because Lemma 1 makes the recursive codewords identical. The exceptional event $c = 0$ has probability 2^{-128} .

The proof format contains twelve categories of opening fields. The implementation exhaustively deconstructs and classifies them, so adding an unclassified field fails compilation. The exact translator is tested against the production additive NTT and wide interleaving. Its functional manifest and entropy table are pinned in `s3_opening_functionals.json` and `s3_minentropy_table.json`.

6 Simulator and concrete ZK bound

Once Theorem 1 and Lemma 1 are separated, the simulator is direct. It receives a public-statement capability containing the setup and digest list, but no witness type. It constructs an unrelated trace, overwrites the public-output region, samples the terminal zerocheck values, and solves the dependent messages in one pass. The patched trace is not an R1CS witness: running the honest prover on it is rejected. After the simulator programs the required fresh oracle points, however, the production verifier accepts the simulated proof. The simulator programs at most 18 points.

The hybrid argument is summarized in Table 2.

Table 2: Game changes and their explicit charges.

Game	Change	Charge
G_0	honest proof	0
G_1	translate the PIOP mask	$720/2^{128}$
G_2	translate μ and g	$1/2^{128}$
G_3	cross the Merkle boundary	$Q_H/2^{16384} + \varepsilon_{\text{RO}}$
G_4	program fresh challenges	$18Q_H/2^{256}$
G_5	solve the terminal messages	$2/2^{128}$

The three algebraic numerators are 720, 1, and 2: the symbolic bad set, $c = 0$, and the two rejected values of the terminal challenge. Thus the part of the bound tracked explicitly by the executable ledger is

$$\varepsilon_{\text{explicit}}(Q_H) = \frac{723}{2^{128}} + \frac{18Q_H}{2^{256}} + \frac{Q_H}{2^{16384}}.$$

At $Q_H = 2^{64}$,

$$-\log_2 \varepsilon_{\text{explicit}}(2^{64}) = 118.502 \dots$$

The term ε_{RO} denotes the ordinary collision and binding term at the framed 256-bit oracle boundary. It is negligible for the stated classical query regime, but it is not included in the quoted 118.502-bit figure.

Theorem 2 (Scoped computational zero knowledge). *Assume the classical programmable random oracle of Assumption 1, the conditioned coverage of Theorem 1, the opening translation of Lemma 1, and freshness of the programmed points. Then both profiles in Table 1 are computationally zero knowledge. For every classical distinguisher making at most Q_H oracle queries,*

$$\text{Adv}_{\text{ZK}} \leq \varepsilon_{\text{explicit}}(Q_H) + \varepsilon_{\text{RO}}.$$

The simulator uses only the public statement, and the verifier is the production verifier.

This theorem is in the classical programmable-ROM. It is not a QROM theorem, and it does not turn the ideal-oracle argument into a standard-model theorem for SHA-256.

7 Extraction and simulation extractability

Zero knowledge says that a proof reveals no witness. Knowledge soundness asks the complementary question: can a successful prover be used to recover one? The two properties use different arguments and have different concrete bounds.

The recording oracle observes the framed leaf queries used to construct the witness commitment. Given a complete valid codeword, the extractor rebuilds the interleaved word, applies the inverse NTT lane by lane, re-encodes the candidate, checks the configured unique-radius condition, and returns the witness half. The system-specific step then reads the 512 message bits of each instance and independently verifies

$$\text{BLAKE3}(x_i) = y_i.$$

Honest recorded commitments yield the true preimages; simulated commitments do not extract.

The executable decoder is an exact-word decoder and candidate validator. It is not a complete Reed–Solomon decoder for every noisy word inside the theoretical radius. The general noisy-word step therefore comes from the Ligerito proximity and unique-decoding argument, not from the current decoder routine.

For simulation extractability, the simulator records

$$(statement\ digest, proof\ nonce, protocol\ version)$$

for each simulated proof. Extraction refuses every recorded prefix and is attempted only on a fresh one. The implemented property is therefore *fresh-prefix weak simulation extractability*. Same-prefix simulation extractability is not claimed.

8 Knowledge-security profiles

The knowledge ledger contains an interactive L0 fold term

$$\kappa_\Gamma = \frac{257}{2^{128}},$$

or 119.994 bits. The classical Fiat–Shamir reduction contributes the factor $Q_H + 1$. At $Q_H = 2^{64}$ it therefore yields

$$-\log_2((Q_H + 1)\kappa_\Gamma) = 55.994 \dots \text{ bits.}$$

This loss is classical; it is not a post-quantum penalty.

The shipped standalone profile must be labelled *100-bit conjectured classical knowledge security*. The theorem column remains 55.994 bits. An unbiased post-commitment beacon removes the Fiat–Shamir query loss but changes the protocol model. A challenge-only $\mathbb{F}_{2^{256}}$ profile would restore a wide standalone margin, but the required machinery does not exist in this implementation. Grinding, a small- Q_H convention, generic parallel repetition, and deleting one fold challenge do not repair the standalone bound.

Table 3: Knowledge claims are separate from the ZK theorem.

Profile	Reduction bound	Status
Standalone FS over $\mathbb{F}_{2^{128}}$	55.994 bits at $Q_H = 2^{64}$	deployment label: 100-bit conjectured classical knowledge security
Unbiased post-commitment beacon over $\mathbb{F}_{2^{128}}$	119.994 bits	100-bit theorem; not standalone
Standalone FS with challenge-only $\mathbb{F}_{2^{256}}$	183.994 bits at $Q_H = 2^{64}$	not implemented

9 Reproduction and limitations

The wire protocol is `flock-zk-fv-v3`; proof I/O version 6 and transcript schema version 5 bind the nonce, oracle frames, and field-mask schema. The two registered circuit digests are:

BLAKE3 batch:

4a398ca2e73d9d1f70611bd6a2a408b0404aabddcf6cf75c29cfa412f72f8423

BLAKE3 fixed digest:

e71bde2b053b1b697412d2d30c2e2e57f7b5f1822484fcbcc15c89ea202890d5

The complete evidence gate is:

```
rustup override set stable
scripts/zk-certify.sh
cargo test --release --workspace --features zk,symbolic
cargo clippy --workspace --all-targets --features zk,symbolic -- -D warnings
cd lean && lake build
```

It replays the symbolic artifacts, transcript schema, oracle framing, production rank gates, simulator, extractor, fixed-digest and batch roundtrips, Lean adapters, knowledge ledger, and the ban on unframed direct SHA-256 calls. Evidence names are compared with script test names in both directions.

The executable evidence is useful because it ties each numerical claim to the shipped code. It is not a substitute for a cryptographic proof or for independent review. The remaining limits are:

- The ZK theorem is classical-ROM only; there is no QROM analysis.
- The standalone knowledge reduction over $\mathbb{F}_{2^{128}}$ gives 55.994 bits at $Q_H = 2^{64}$, not a 100-bit theorem.
- The executable decoder validates candidates but does not implement general arbitrary-error Reed–Solomon decoding.
- Only the two exact registry entries are enabled. Other Flock statement families remain non-ZK.
- Side-channel resistance, recursive composition, and same-prefix simulation extractability are out of scope.
- The construction and proof have not received independent cryptographic review.

Within these limits, the conclusion is precise: for the registered BLAKE3 batch and fixed-digest relations, Flock now has a public-statement-only simulator whose proofs are accepted by the production verifier, together with an explicit distribution argument for every witness-dependent part of the transcript. That is the sense in which these Flock profiles are zero knowledge.

References

- [1] E. Ben-Sasson, A. Chiesa, and N. Spooner. *Interactive Oracle Proofs*. TCC 2016; IACR ePrint 2016/116.
- [2] A. Chiesa, M. A. Forbes, and N. Spooner. *A Zero Knowledge Sumcheck and its Applications*. IACR ePrint 2017/305.
- [3] A. Novakovic and G. Angeris. *Ligerito: A Small and Concretely Fast Polynomial Commitment Scheme*. IACR ePrint 2025/1187.
- [4] B. E. Diamond and J. Posen. *Polylogarithmic Proofs for Multilinears over Binary Towers*. IACR ePrint 2024/504.
- [5] J. T. Schwartz. *Fast Probabilistic Algorithms for Verification of Polynomial Identities*. Journal of the ACM 27(4), 701–717, 1980.
- [6] R. Zippel. *Probabilistic Algorithms for Sparse Polynomials*. EUROSAM 1979, LNCS 72, 216–226.
- [7] Succinct Labs. *Flock*. <https://github.com/succinctlabs/flock>.